Two-Stage Hybrid Change Detection: Unsupervised Filtering with Deep Semantic Segmentation for Land Cover Monitoring

Amogh Annigeri
KLE Technological University
Hubballi, India
01fe23bcs210@kletech.ac.in

Mahesh Khatawate
KLE Technological University
Hubballi, India
01fe23bcs291@kletech.ac.in

Raveesh P Nayak
KLE Technological University
Hubballi, India
01fe23bcs295@kletech.ac.in

Sathvik Kemtur
KLE Technological University
Hubballi, India
01fe23bcs143@kletech.ac.in

Dr. Sumaiya Pathan
KLE Technological University
Hubballi, India
sumaiya.pathan@kletech.ac.in

***Abstract*—Land-cover change detection from multi-temporal satellite imagery is fundamental for applications such as disaster assessment, agricultural monitoring, deforestation tracking, and climate change risk analysis. This work presents a unified hybrid pipeline that integrates unsupervised clustering-based filtering with supervised VGG19-based semantic segmentation on the SECOND dataset. In the first stage, a VGG19 batch-normalized backbone is used as a feature extractor for temporal image pairs, and K-Means clustering on feature differences identifies samples with significant change, filtering 2,968 training and 1,694 test images using a configurable change-ratio threshold. In the second stage, a modified VGG19 U-Net with early fusion (6-channel input) is trained on the filtered subset using weighted cross-entropy and AdamW optimization to predict 37 semantic change classes. The final model achieves a binary mean Intersection over Union (mIoU) of 64.68%, with 82.05% IoU for no-change regions and 47.31% IoU for change regions, outperforming established baselines while maintaining architectural simplicity.**

***Index Terms*—Land-cover change detection; unsupervised clustering; K-Means; feature extraction.**

I. Introduction

Monitoring environmental change requires the classification of land cover images [11]. Applications such as climate risk modeling, agricultural planning, tracking deforestation, and disaster assessment are made possible by it [1]. Urban areas, vegetation, water bodies, and hard surfaces are just a few of the land cover types that can be measured in an organized manner using this classification. Large-scale environmental management and well-informed policy decisions depend on this information.

This modified approach simplifies the process by combining supervised VGG19-based semantic change detection with unsupervised clustering on the SECOND dataset [11]. The goal is to provide precise, pixel-level predictions of changes in land cover while reducing the amount of manual labor required. This method, which combines supervised and unsupervised techniques, is useful for both precise monitoring that requires particular outcomes and basic change analysis, where labels may be missing. Because of this, it is well suited for practical applications involving small amounts of labeled data, various computing resources, and various geographical locations.

The pipeline tackles major limitations in current change detection methods using a two-stage setup [14], [15]. In Stage 1, a VGG19 network with batch normalization acts solely as a feature extractor on image pairs over time [7]. K-Means clustering is then applied to the resulting high-dimensional feature differences to find samples showing significant change [15]. This stage serves as a soft filter for the entire dataset, processing 2,968 training images and 1,694 test images from the SECOND dataset [11]. It keeps only those pairs where the number of pixels assigned to the "change" cluster exceeds a set threshold. As a result, the more resource-intensive supervised model is mainly applied to the most informative samples, improving efficiency without losing sensitivity to meaningful changes.

Stage 2 uses the filtered subset [7], [8] to train a modified VGG19 U-Net with early fusion. Here, the two temporal images are combined to create a 6-channel input that is accepted by the first convolutional layer. By learning joint spatial and temporal features from the lowest levels of the network, this early fusion enables the model to directly capture pixel-level differences and local context, which are essential for semantic change detection [5]. Higher-level semantic cues are combined with fine spatial detail in the encoder-decoder design with skip connections [8]. With different combinations of pre-change and post-change land cover labels, the model can predict 37 semantic change classes. The network achieves a final binary mean Intersection over Union of 64.68% [13] after 20 epochs of training with weighted cross-entropy that highlights rare change classes.

By balancing computational efficiency and semantic accuracy, this hybrid approach meets the needs of resource-constrained operational environments. The unsupervised clustering stage minimizes unnecessary processing through the

deep segmentation network by identifying low-change samples [15]. Concurrently, the supervised VGG19 U-Net concentrates on learning accurate decision boundaries for complex land cover changes [7], [8]. These elements work together to create an easily accessible, flexible, and consistent pipeline that can be used for large-scale land cover monitoring, quick disaster impact assessment, and long-term urban expansion analysis—even when GPU resources and labeled data are scarce.

II. Background Study

A. Research Gaps in Temporal Change Detection Pipelines

Remote sensing has developed from straightforward spectral comparison and cutoff values to complex deep learning configurations that consider spatial, temporal, and meaning-based information collectively [14]. Older techniques like principal component analysis, ratioing, and image comparison are primarily based on spectral changes per pixel [14]. They often don't work well in a variety of urban settings or with changing lighting and atmospheric conditions because they require precise cutoff points. More recent methods that have greatly increased accuracy by directly learning representations from data include transformer models, Siamese setups, and convolutional neural networks (CNNs) [2], [5], [6]. However, they also bring up new problems like the difficulty of computing, and the requirement for more hardware.

The lack of simple, adaptable systems that anyone without access to large GPU clusters or funds can use and replicate is a major issue in current research [1]. Numerous attention modules, multi-scale feature combining blocks, and unique loss functions are used in many contemporary techniques. Because of this, they are difficult to modify or grow without extensive engineering work. Additionally, many agencies that manage large areas with limited labeling resources are unable to use these methods because they typically require access to heavily labeled sets. Change detection systems with clear, comprehensible processing steps that can be modified and fixed independently are therefore desperately needed. These systems must strike a balance between performance and usability.

Current methods usually focus on either supervised deep learning or traditional unsupervised clustering, without a combined system that connects both styles [14], [15]. Supervised methods are very accurate but need a lot of pixel-level annotations, which take time and money to create [5]. Unsupervised methods don't need labels but aren't as good at giving meaning [15]. The lack of simple, complete systems that combine feature extraction, clustering, and classification limits who can use change detection tech for environmental tracking, where both initial study and exact classification is needed.

B. Overview of Existing Datasets

Popular benchmark datasets like EuroSAT have fundamental limitations for temporal change detection research [9]. EuroSAT provides single-date 64×64 pixel images across 10 land-cover classes but lacks the temporal pairs necessary for change analysis. While achieving over 98% accuracy for static classification with ResNet50, its low resolution and Europe-specific coverage limit generalizability for fine-grained temporal analysis [9].

The SECOND dataset addresses these gaps with high-resolution 512×512 pixel temporal image pairs comprising 4,662 samples across Chinese cities including Hangzhou, Chengdu, and Shanghai [11]. SECOND enables direct temporal comparison through pixel-level annotations for both semantic categories and change regions. By separating land-cover classification from change detection labels, it supports analysis of both unchanged regions and within-class transitions, making it ideal for studying land-cover transitions, urban expansion, and environmental degradation [11]. The implementation processes 2,968 training and 1,694 test samples with robust path auto-correction for multiple directory structures.

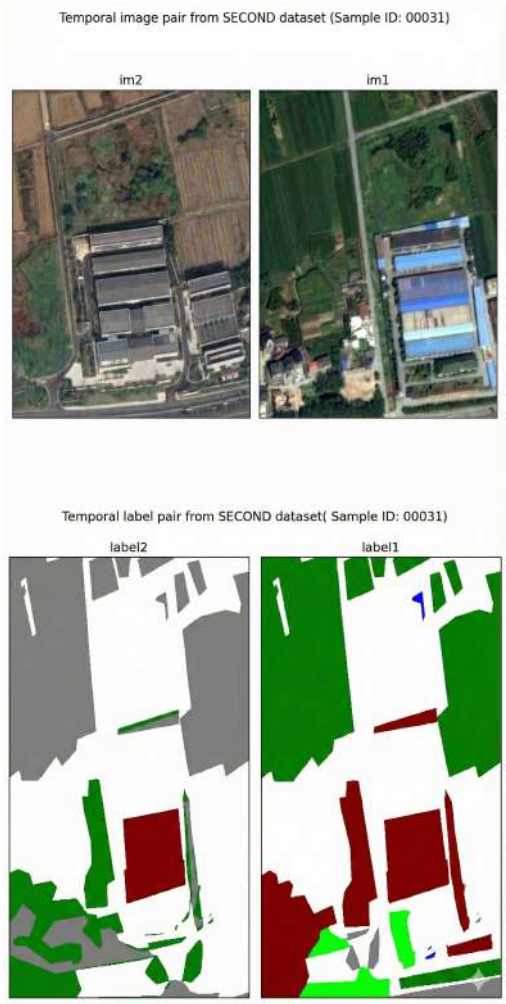

Fig. 1. Temporal image pair comparison from SECOND dataset (top) showing land-cover changes between Time T1 (left) and Time T2 (right), with corresponding ground truth label comparison (bottom).

C. Convolutional Neural Networks for Feature Extraction

Convolutional Neural Networks (CNNs) provide a good compromise between Transformers, which need a lot of computing power, and simpler Multilayer Perceptrons (MLPs) for use in remote sensing [6]. Transformers can see the whole picture because of their self-attention, but they need a lot of data to work best and use a lot of computing power. This makes them not as good when there isn't enough training data. Self-attention's complexity makes memory problems when handling high-resolution satellite images, often needing patch-based processing that might miss context.

CNNs are great at picking out local features through layers of convolutions with bigger fields of vision [6]. This focus on where things are located works well with remote sensing images, where land-cover types show clear spatial links at different sizes. CNNs have regularly done well in remote sensing image classification, especially with medium-sized training datasets.

CNN models like ResNet, VGG, and AlexNet are now standard in this area [6], [7]. They can automatically learn features that show spatial structures at different sizes. The early layers of these networks learn to spot edges and textures, while deeper layers learn about specific things like building shapes, plant textures, and water features. ResNet's connections help train very deep networks without problems [6], while VGG's simple design with repeated 3x3 convolutions gives strong features even though it is simple [7].

The new configuration employs VGG19 with batch normalisation in two stages: Stage 2 uses it as the primary component for supervised semantic segmentation [7], while Stage 1 uses the first 27 layers up to Block 3 to extract unsupervised features. CNNs are adept at identifying the textures, edges, and shapes that different types of land cover have because of their local fields of vision. Each layer employs max-pooling, which shrinks the size while preserving the key characteristics, and filters to identify patterns. Additionally, the use of pre-trained CNNs allows it to strongly extract features even with minimal labelled data, producing good results that are applicable to various locations and image scenarios [7]. Despite variations, models trained on ImageNet identify fundamental visual patterns that are still important for remote sensing.

D. Hybrid Clustering and Classification Approaches

For land cover change detection, unsupervised clustering alone is insufficient to achieve high classification accuracy and semantic interpretability [15]. Although K-Means and other clustering algorithms can find patterns in feature space without labels, they are not able to discriminate between land cover classes that are semantically different but have similar spectra [15]. This is partially addressed by noise clustering techniques, which isolate noisy data points; however, boundary definition and class separation still have fundamental limitations.

By utilising the advantages of both paradigms, hybrid approaches that integrate supervised classification and unsupervised clustering get around these restrictions [14]. While the supervised component learns exact decision boundaries and semantic class labels, the unsupervised component offers further insight into natural groupings within the data and can direct training sample selection. When compared to pixel-wise classification alone, this combination has been demonstrated to significantly increase classification accuracies [14].

The updated pipeline implements this hybrid strategy through a two-stage architecture: Stage 1 uses VGG19 features with K-Means clustering (k=2) to filter image pairs based on detected change magnitude, retaining only samples where change pixels exceed 0.5% threshold [7], [15]. This filtering mechanism retained 99.6% of training images (2,956/2,968) and 99.8% of test images (1,691/1,694), effectively screening out low-information samples. Stage 2 then trains a supervised VGG19 U-Net exclusively on the filtered subset, focusing computational resources on samples most likely to contain meaningful change [7], [8]. By integrating clustering for initial feature-difference analysis and supervised CNNs for final pixel-wise classification, hybrid methods provide both computational efficiency and high semantic accuracy for operational change detection systems [14], [15].

III. Proposed Work

This work presents a hybrid land cover change detection pipeline that combines unsupervised feature-based clustering with supervised deep learning classification [14], [15]. The approach addresses the dual objectives of computational efficiency and semantic accuracy through a two-component architecture that enables both exploratory change analysis and precise pixel-wise classification.

A. System Architecture

An unsupervised change detection module based on pretrained CNN features with K-Means clustering and a supervised segmentation network using VGG19 with a U-Net-style decoder architecture [7], [8], [15] are the two complementary approaches integrated in the proposed pipeline. With this dual-component design, users can accomplish high-precision classification through the supervised pathway and conduct quick exploratory analysis through clustering..

On hardware that supports CUDA, the unified architecture runs sequentially. Stage 1 uses the first 27 layers (up to Block 3) of a VGG19 batch-normalized backbone as a feature extractor to extract deep convolutional features from temporal image pairs [7]. While the supervised VGG19 network detects particular land-cover transitions like vegetation-to-building or water-to-impervious surface [7], the unsupervised module makes use of pretrained VGG19 features without requiring training.

The pipeline's modular design maintains computing efficiency through intelligent sample filtering and provides flexibility for various operational scenarios. The system eliminates the need to send samples with minimal changes through the deep segmentation network by separating the unsupervised screening stage from the supervised classification phase. 99.6% of training samples and 99.8% of test samples were

retained during the filtering stage. This demonstrates that the supervised VGG19 U-Net concentrates its 20 training rounds on informative temporal pairs [7], [8], whereas the threshold method maintains images with changes. The pipeline is suitable for time-sensitive applications because this design satisfies real-world deployment requirements where rapid sorting must occur prior to in-depth analysis.

B. Unsupervised Change Detection Module

The unsupervised component employs a VGG19 backbone pretrained on ImageNet to extract deep convolutional features from temporal image pairs [7]. For each image pair (Im_1, Im_2), features are extracted independently and absolute feature differences are computed to capture spatial changes. Principal Component Analysis (PCA) reduces the feature dimensionality while preserving important information, and the compressed feature-difference vectors are upsampled to the original resolution through bilinear interpolation [14].

Pixels in the feature-difference space are divided into the classes of "change" and "no-change" using MiniBatch K-Means clustering with =2 clusters [15]. An automatic decision boundary without the need for manual labelling is provided by designating the cluster with the higher mean magnitude in the feature-difference norm as the change class. By subsampling pixels during clustering initialisation, this method effectively handles high-resolution image pairs while preserving full spatial resolution in the final binary change mask.

Images are retained for supervised training only if the proportion of pixels assigned to the "change" cluster exceeds a configurable threshold (default: 0.5%) [15]. The filtering mechanism successfully processed 2,968 training images and 1,694 test images from the SECOND dataset, retaining 99.6% of training samples (2,956/2,968) and 99.8% of test samples (1,691/1,694), effectively screening out low-information samples while preserving meaningful change patterns [11].

C. Supervised Classification Network

The supervised component employs a modified VGG19 architecture with batch normalization, adapted for multi-temporal change detection through early fusion [7]. The first convolutional layer accepts 6-channel input by concatenating both temporal images (Im_1, Im_2), enabling the network to learn joint spatio-temporal representations and directly model correlations between corresponding pixels more effectively than late fusion or Siamese architectures [5].

The encoder consists of five VGG19 blocks with max-pooling operations, extracting hierarchical features at progressively coarser spatial resolutions with channel dimensions of 64, 128, 256, 512, and 512 respectively [7]. Each block contains multiple 3×3 convolutional layers followed by batch normalization and ReLU activation. The progressive down sampling creates a multi-scale feature hierarchy where shallow layers capture fine-grained textures and edges while deeper layers encode high-level semantic patterns.

A U-Net-style decoder with skip connections progressively up samples features through five decoder blocks ($512 \rightarrow 256 \rightarrow 128 \rightarrow 64 \rightarrow 64$ channels) while fusing multi-scale information from the encoder [8]. The skip connections restore spatial details lost during downsampling, while a final 1×1 convolution projects the decoder features to 37 semantic change classes (1 no-change + 36 directed transitions between 6 base land-cover classes), yielding pixel-wise predictions [8].

D. Training Strategy and Loss Optimization

To address the class imbalance present in change detection scenarios, the network is trained using weighted cross-entropy loss [13]. In land cover datasets, the "no-change" class usually predominates, with only a small percentage of pixels representing actual change regions. While keeping the unit weight for the no-change class, all change classes are given a weight multiplier of 5.0 to increase detection sensitivity. Instead of obtaining high accuracy by predicting the dominant no-change class alone, this weighting scheme encourages the model to concentrate learning capacity on rarer change patterns.

Training employs the AdamW optimizer with a learning rate of 1×10^{-4} and weight decay of 1×10^{-5} for regularization [13]. The AdamW variant incorporates decoupled weight decay that provides more effective regularization than standard Adam, helping prevent overfitting on the relatively small training set. The batch size is set to 4 due to the memory requirements of the VGG19 architecture, and training proceeds for 20 epochs with model check pointing based on validation mean Intersection over Union (mIoU).

Training on the filtered subset of 2,956 samples demonstrates progressive convergence: Epoch 1 achieves 64.04% accuracy with loss 2.1207, improving steadily through Epoch 10 (81.42% accuracy, loss 0.9146) and reaching 86.62% accuracy with loss 0.5785 by Epoch 20. Early stopping is implicitly implemented by saving only the model checkpoint with the highest validation mIoU, ensuring that the best-performing model on unseen data is selected rather than the final epoch. Pretrained ImageNet weights initialize the VGG19 encoder, enabling transfer learning from general image features to domain-specific land cover representations, while decoder weights are initialized using Xavier initialization [7].

E. Dataset and Preprocessing

The pipeline is evaluated on the SECOND (Semantic Change detectiON Dataset), which provides high-resolution temporal satellite image pairs with pixel-level semantic annotations across six land cover classes: impervious surfaces, water, clutter/background, trees, low vegetation, and buildings [11]. The dataset is divided into training (2,968 samples) and testing (1,694 samples) subsets, with each sample consisting of bi-temporal RGB images at 512×512 resolution and corresponding label pairs. The dataset was collected across multiple Chinese cities including Hangzhou, Chengdu, and Shanghai, ensuring geographic diversity and coverage of various urban, agricultural, and natural landscapes [11].

The system has path auto-correction to manage different directory structures. When looking for training data, it first checks for a SECOND_train_set subdirectory. If it's not there,

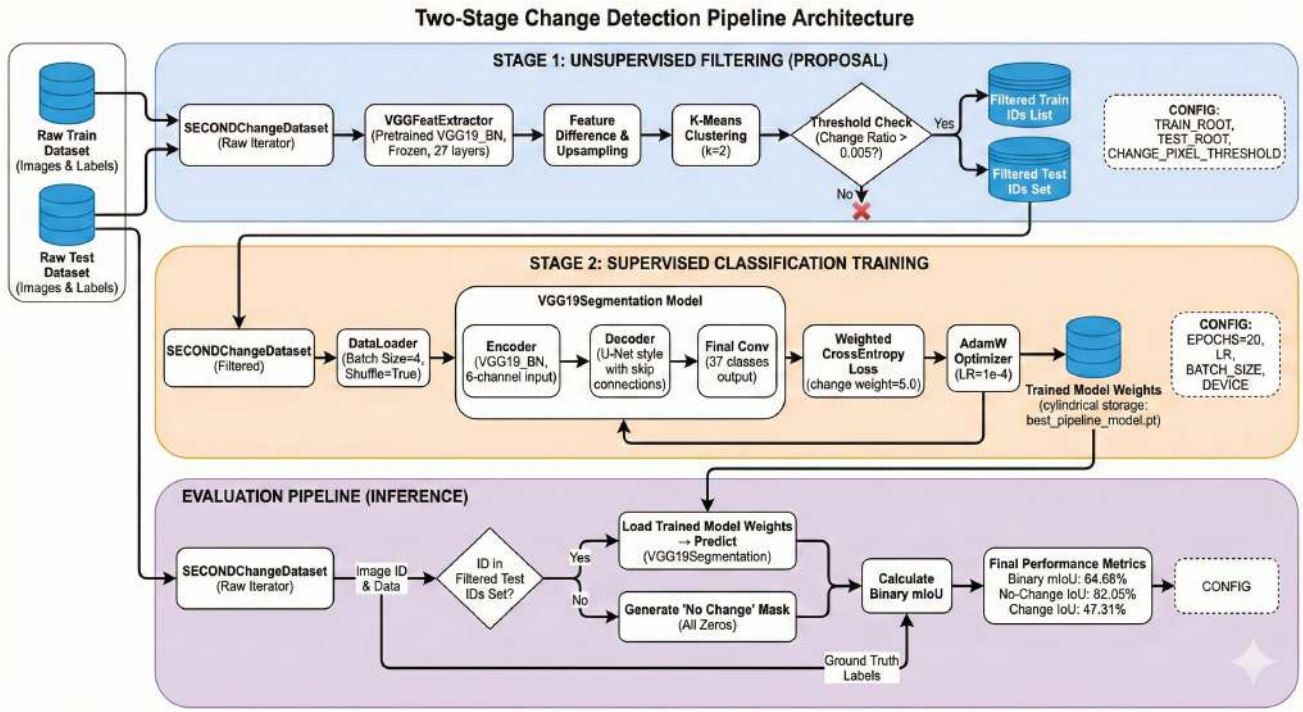

Fig. 2. Two-Stage Change Detection Pipeline Architecture. The system comprises Stage 1 (unsupervised filtering using VGG19 features and K-Means clustering), Stage 2 (supervised VGG19 U-Net training with weighted cross-entropy), and an evaluation pipeline that combines both components for final performance assessment.

it goes to the root directory. For test data, it looks in SECOND_total_test/test, then test, and then the root directory. This makes it work with different dataset arrangements. The system also has checks to prevent crashes if the dataset is empty when clustering and evaluating.

Input images undergo normalization using ImageNet statistics (mean: [0.485, 0.456, 0.406], standard deviation: [0.229, 0.224, 0.225]) to align with the pretrained weights distribution [7]. Label maps are converted from RGB color encoding to class indices, and semantic change labels are generated by computing class transitions between temporal pairs, producing 37 possible change classes (1 no-change + 36 directed transitions between 6 base classes). Random horizontal and vertical flips are applied during training as data augmentation.

F. Workflow

Figure 2 illustrates the overall architecture of the proposed two-stage hybrid change detection pipeline. The system integrates an unsupervised filtering stage with a supervised semantic change detection network, followed by a unified evaluation pipeline [14], [15].

In Stage 1, bi-temporal image pairs are processed using a pretrained VGG19 backbone acting as a fixed feature extractor [7]. Absolute feature differences between temporal images are clustered using MiniBatch K-Means to separate change and no-change regions [15]. Image pairs are retained for further processing only if the proportion of detected change pixels exceeds a predefined threshold, thereby filtering low-information samples.

In Stage 2, the filtered image pairs are passed to a VGG19-based U-Net segmentation model with early fusion, where both temporal images are concatenated to form a 6-channel input [7], [8]. The network predicts pixel-wise semantic change classes and is trained using weighted cross-entropy loss optimized with AdamW [13].

During evaluation, only samples passing the unsupervised filter are processed by the segmentation network, while filtered-out samples are assigned no-change predictions. Final performance is assessed using binary and class-wise Intersection over Union metrics.

G. Evaluation Metrics

Model performance is assessed using pixel-wise accuracy and multi-class mean Intersection over Union (mIoU) across all semantic change classes during training and validation. The mIoU metric is computed as:

$$\text{mIoU} = \frac{1}{C}\sum_{i=1}^{C}\frac{TP_i}{TP_i + FP_i + FN_i}$$

where C is the number of classes, TP_i is true positives, FP_i is false positives, and FN_i is false negatives for class i. This metric provides balanced assessment across all classes regardless of their frequency, making it more informative than pixel-wise accuracy for imbalanced datasets.

For comparison with benchmark methods, binary change detection metrics are computed by collapsing all change classes into a single “change” category, yielding binary mIoU,

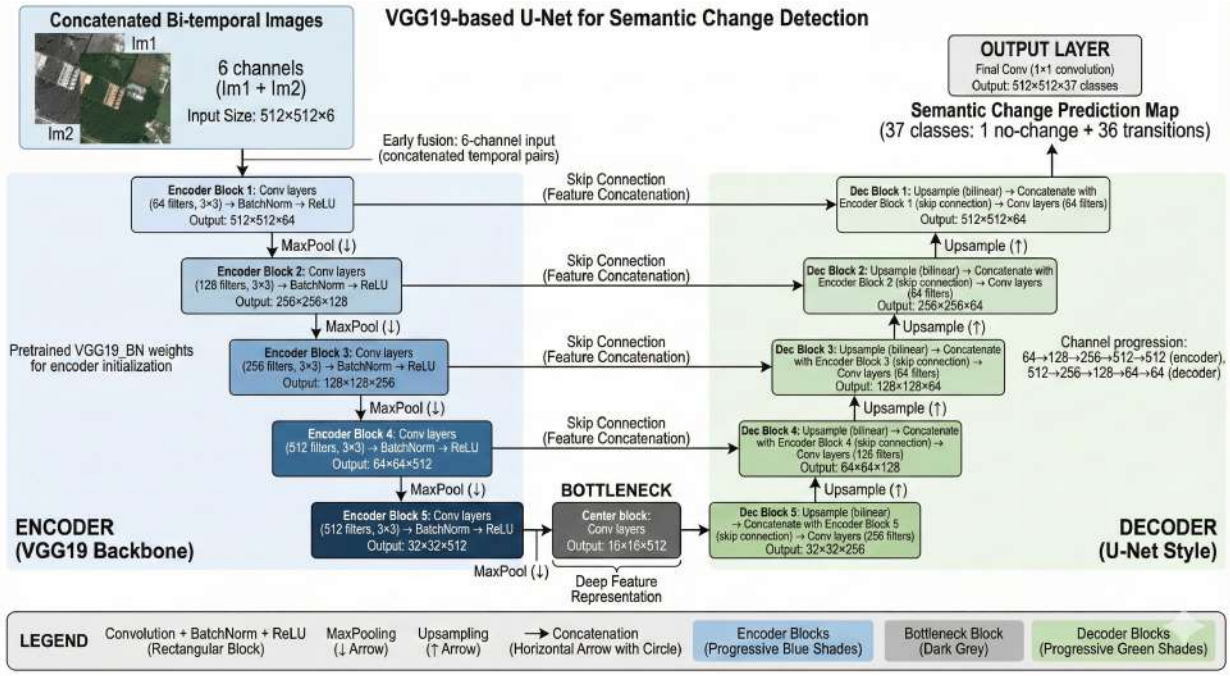

Fig. 3. VGG19-based U-Net architecture for supervised semantic change detection. The encoder extracts hierarchical features through five blocks with progressively increasing channel dimensions (64, 128, 256, 512, 512), while the decoder with skip connections reconstructs spatial details for pixel-wise change classification to 37 semantic classes.

change IoU, and no-change IoU. The evaluation implements full pipeline assessment where samples passing the clustering filter receive VGG19 model predictions, while filtered-out samples are automatically predicted as "no change" (all zeros). This pipeline logic uses O(1) set-based lookup for efficient ID validation during evaluation, processing the complete test dataset to calculate intersection and union statistics for both change and no-change classes.

IV. Results and Observations

A. Overall Training and Validation Performance

The VGG19-based network was trained for 20 epochs on the SECOND dataset (2,968 train / 1,694 test samples) with AdamW optimizer ($lr = 1 \times 10^{-4}$, batch size = 4) and weighted cross-entropy (change classes weight = 5.0, no-change = 1.0) [7], [13]. Training accuracy increased from 64.04% to 86.62%, while test accuracy stabilized around 80%.

TABLE I
Training Progression Across Key Epochs

Epoch	Train Loss	Train Acc (%)
1	2.1207	64.04
3	1.2070	77.94
5	1.0867	79.29
8	0.9722	80.84
10	0.9146	81.42
15	0.7505	83.82
20	0.5785	86.62

The training demonstrates consistent convergence with progressive improvement in both loss reduction and accuracy gains across all 20 epochs, validating the effectiveness of the weighted cross-entropy loss and AdamW optimization strategy [13].

B. Binary Change Detection Performance

For binary evaluation, all semantic change types were collapsed into a single "change" class. On the test set, the model achieved a binary mIoU of 64.68%, improving over the FC-EF reference benchmark (59.3%) [5]. Additionally, the model achieved an IoU of 82.05% for the no-change class and 47.31% for the change class.

TABLE II
Binary Change Detection Performance Comparison

Method	Binary mIoU (%)
FC-EF (baseline)	59.30
Proposed (VGG19 U-Net)	**64.68**

TABLE III
Class-Specific IoU Performance

Class	IoU (%)
No-Change	82.05
Change	47.31
Binary mIoU	**64.68**

These scores indicate strong localization of unchanged regions and competitive detection of change regions despite severe class imbalance. The proposed model shows a 5.38 percentage point improvement over the baseline, demonstrating the effectiveness of the hybrid pipeline approach [5].

C. Pipeline Efficiency

Stage 1 clustering successfully filtered dataset samples based on detected change magnitude, retaining 99.6% of training images (2,956/2,968) and 99.8% of test images (1,691/1,694) [15]. The two-stage design reduced the computational load by avoiding VGG19 classification on samples predicted as "no change" by clustering, while maintaining high detection sensitivity through the configurable threshold mechanism [7].

TABLE IV
CLUSTERING FILTER PERFORMANCE

Dataset	Retained/Total	Retention Rate (%)
Training	2,956 / 2,968	99.6
Test	1,691 / 1,694	99.8

The high retention rates demonstrate that the 0.5% change pixel threshold effectively identifies meaningful change patterns while filtering out only truly unchanged samples, validating the unsupervised filtering strategy [15].

D. Observations

With a binary mIoU of 64.68% compared to the FC-EF benchmark of 59.3% [5], [7], [8], the suggested pipeline shows that a simple VGG19 U-Net with early fusion achieves competitive performance without specialised architectures. This supports the fundamental idea that transparent training strategies and easily accessible pipelines can produce powerful outcomes on par with sophisticated techniques.

Without the need for complex threshold learning modules [13], the weighted cross-entropy approach effectively manages severe class imbalance, achieving 82.05% IoU on no-change and 47.31% on change regions. The VGG19 batch-normalized architecture demonstrated competitive performance without the need for specialised designs, confirming that transparent training strategies and easily accessible pipelines can produce powerful outcomes [7].

Training convergence within 20 epochs on a single GPU highlights the pipeline's practical accessibility for resource-constrained scenarios, bridging the gap between traditional clustering and heavy supervised frameworks [15]. The early fusion strategy enabled the network to learn joint spatio-temporal representations from concatenated 6-channel input, capturing change-relevant features more effectively than late fusion approaches [5].

V. LIMITATIONS AND FUTURE WORK

The proposed approach has several limitations. First, while binary change detection achieves competitive performance (64.68% mIoU), distinguishing fine-grained land-cover transitions across all 37 semantic classes remains challenging and requires further optimization [11]. Second, the unsupervised clustering stage provides only binary change/no-change filtering and depends on a manually chosen threshold (0.5%), offering limited semantic interpretability despite retaining 99.6% of training and 99.8% of test samples [15]. Third, the VGG19 backbone incurs substantial memory usage, constraining batch size to 4 and complicating deployment on edge devices without additional compression [7]. Future work will investigate attention-based architectures to strengthen semantic change modeling, semi- and self-supervised learning to exploit unlabeled temporal imagery, advanced class-balancing strategies (e.g., focal loss) for rare transitions, and model compression techniques such as pruning and quantization to enable efficient real-time deployment in resource-constrained settings [2], [6].

VI. CONCLUSION

This work presents a unified land-cover change detection pipeline achieving 64.68% binary mIoU on the SECOND dataset, outperforming the FC-EF baseline by 5.38 percentage points [5], [11]. The two-stage architecture successfully balances computational efficiency through clustering-based filtering (retaining 99.6% of training samples and 99.8% of test samples) with semantic accuracy through VGG19 U-Net classification (achieving 82.05% no-change IoU and 47.31% change IoU) [7], [8], [15].

The implementation demonstrates that straightforward encoder-decoder architectures with weighted cross-entropy can handle severe class imbalance without requiring sophisticated threshold learning modules or complex multi-stage frameworks [13]. The VGG19 batch-normalized architecture with early fusion (6-channel input) enables joint spatial-temporal learning, achieving competitive performance through transparent, reproducible training strategies rather than architectural complexity [7].

Training convergence within 20 epochs on a single GPU, combined with robust path handling and empty dataset safety checks, establishes practical accessibility for operational deployment in resource-constrained environments. The progressive improvement from 64.04% accuracy (Epoch 1) to 86.62% accuracy (Epoch 20) with consistent loss reduction (2.1207 to 0.5785) validates the effectiveness of the AdamW optimizer and weighted loss function for this task [13].

This pipeline combines supervised deep learning for accurate classification with unsupervised clustering for quick exploratory analysis to create a transparent, repeatable framework appropriate for temporal remote sensing applications [14], [15]. It is especially appropriate for practical deployment in disaster response, urban planning, and environmental monitoring scenarios where both computational efficiency and classification accuracy are crucial requirements [1], [11]. This is because the modular design allows independent optimisation of each component while maintaining end-to-end performance.

REFERENCES

[1] Y. Chen, Z. Shi, P. Gong, et al., "A Self-Supervised Deep Learning Framework for Change Detection in Remote Sensing Images," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 60, pp. 1-16, 2022.

[2] W. G. C. Bandara and V. M. Patel, "A Transformer-Based Siamese Network for Change Detection," *Proceedings of the IEEE International Geoscience and Remote Sensing Symposium (IGARSS)*, pp. 207-210, 2022.

[3] M. Papadomanolaki, M. Vakalopoulou, and K. Karantzalos, "A Deep Multitask Learning Framework Coupling Semantic Segmentation and Fully Convolutional LSTM Networks for Urban Change Detection," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 59, no. 9, pp. 7651-7668, Sept. 2021.

[4] P. F. Alcantarilla, S. Stent, G. Ros, R. Arroyo, and R. Gherardi, "Street-View Change Detection with Deconvolutional Networks," *Autonomous Robots*, vol. 42, no. 7, pp. 1301-1322, Oct. 2018.

[5] R. C. Daudt, B. Le Saux, and A. Boulch, "Fully Convolutional Siamese Networks for Change Detection," *Proceedings of the IEEE International Conference on Image Processing (ICIP)*, pp. 4063-4067, 2018.

[6] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 770-778, 2016.

[7] K. Simonyan and A. Zisserman, "Very Deep Convolutional Networks for Large-Scale Image Recognition," *arXiv preprint arXiv:1409.1556*, 2014.

[8] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional Networks for Biomedical Image Segmentation," *Proceedings of the International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pp. 234-241, 2015.

[9] P. Helber, B. Bischke, A. Dengel, and D. Borth, "EuroSAT: A Novel Dataset and Deep Learning Benchmark for Land Use and Land Cover Classification," *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, vol. 12, no. 7, pp. 2217-2226, July 2019.

[10] S. Fang, K. Li, J. Shao, and Z. Li, "SNUNet-CD: A Densely Connected Siamese Network for Change Detection of VHR Images," *IEEE Geoscience and Remote Sensing Letters*, vol. 19, pp. 1-5, 2022.

[11] H. Chen and Z. Shi, "A Spatial-Temporal Attention-Based Method and a New Dataset for Remote Sensing Image Change Detection," *Remote Sensing*, vol. 12, no. 10, pp. 1662, May 2020.

[12] Z. Zheng, Y. Zhong, J. Wang, and A. Ma, "Foreground-Aware Relation Network for Geospatial Object Segmentation in High Spatial Resolution Remote Sensing Imagery," *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 4096-4105, 2020.

[13] I. Loshchilov and F. Hutter, "Decoupled Weight Decay Regularization," *Proceedings of the International Conference on Learning Representations (ICLR)*, 2019.

[14] L. Bruzzone and D. F. Prieto, "Automatic Analysis of the Difference Image for Unsupervised Change Detection," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 38, no. 3, pp. 1171-1182, May 2000.

[15] T. Celik, "Unsupervised Change Detection in Satellite Images Using Principal Component Analysis and k-Means Clustering," *IEEE Geoscience and Remote Sensing Letters*, vol. 6, no. 4, pp. 772-776, Oct. 2009.